

# Z6110X0035: Introduction to Cloud Computing - Economics of IaaS Cloud Providers

Lecturer: Prof. Zichen Xu

# Outline

## Background

## Motivations, Research Problem and Thesis Focus

## Solutions

### Federation of Clouds

- Resource Provisioning Policies to Increase Profit

- Financial Option Market Model

### Single Cloud Provider

- Revenue Management with Optimal Capacity Control

- An Auction Mechanism for a Cloud Spot Market

- Spot instance pricing as a Service

## Conclusions and Future Directions

# Outline

## Background

Motivations, Research Problem and Thesis Focus  
Solutions

- Federation of Clouds

  - Resource Provisioning Policies to Increase Profit

  - Financial Option Market Model

- Single Cloud Provider

  - Revenue Management with Optimal Capacity Control

  - An Auction Mechanism for a Cloud Spot Market

  - Spot instance pricing as a Service

Conclusions and Future Directions

# Cloud Computing

Allowing businesses to outsource their IT facilities to cloud providers

Avoid expensive up-front investments of establishing their own infrastructure

Long-held dream of computing as a utility

On-demand delivery of IT services

Customers pay for what they use

Virtualized resources

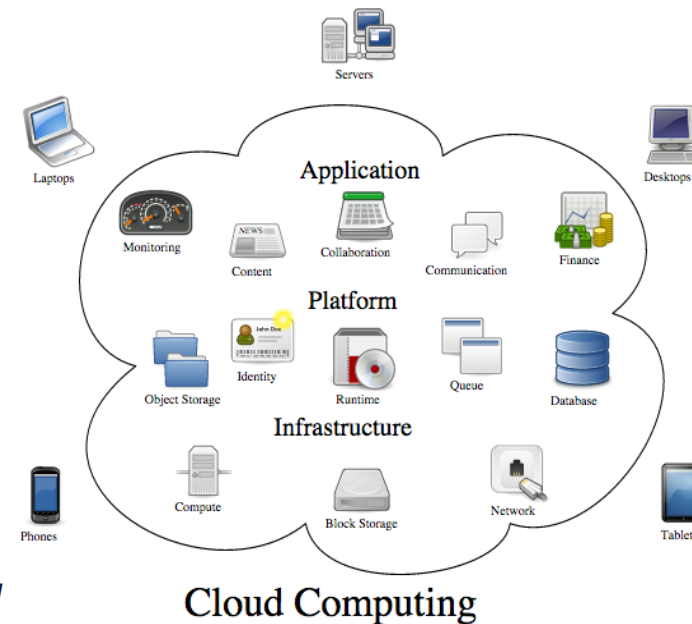


Figure: [http://en.wikipedia.org/wiki/Cloud\\_computing](http://en.wikipedia.org/wiki/Cloud_computing)

# Types of Service Models

## Software as a Service (SaaS)

It provides applications and software to the customer in utility-based model

Accessible from a thin client interface such as a Web browser

Example: Salesforce.com

Customer relationship management (CRM) as a service.

## Platform as a Service (PaaS)

It provides programming languages and tools to  
deploy application onto the cloud infrastructure

Example: Google App Engine

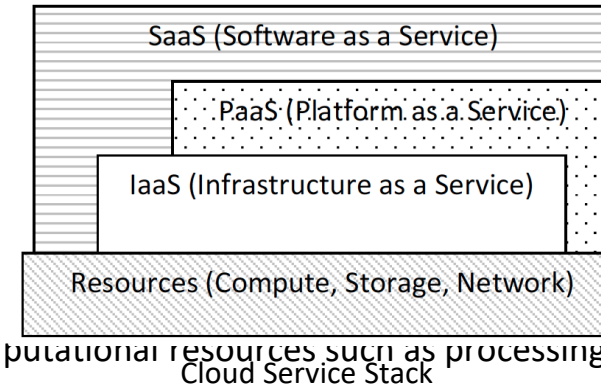
facilities to build an reliable and scalable application

## Infrastructure as a Service (IaaS)

It provides capabilities for the customers to provision computational resources such as processing, storage, network, and other fundamental computing resources

Virtual Machines (VMs)

Example: Amazon EC2/S3



# Outline

Background

**Motivations, Research Problem and Thesis Focus**

Solutions

- Federation of Clouds

  - Resource Provisioning Policies to Increase Profit

  - Financial Option Market Model

- Single Cloud Provider


  - Revenue Management with Optimal Capacity Control

  - An Auction Mechanism for a Cloud Spot Market

  - Spot instance pricing as a Service

Conclusions and Future Directions

# International Data Corporation (IDC) forecast



Worldwide public IT cloud services spending in billion dollars by service model

# Motivation (1)

The primary focus of designers in traditional distributed system:

Improving system performance in terms of response time, throughput system, and utilization;

Extensively studied in Cluster and Grid Computing in the past 15 years.

**Cloud Computing and Computational Economy**

New objective functions: price, cost, and budget

 **Money**

This has shifted a traditional distributed system into a two-party computation:

Cloud providers and Cloud customers

Pricing and cost as the bridge



# Motivation (2)

There is a large body of research devoted to

- Minimizing cost

- Cloud customers

Relatively less work has been done

- Maximizing profit

- The provider's side

Techniques and mechanisms that help providers to maximize their profit while still selling their services competitively.

# Research Problem

This thesis tackles the research challenges arising from the following topic:

*" Designing market and economics-inspired algorithms and resource allocation mechanisms to maximize IaaS cloud providers' profit honoring the QoS constraints associated w SLA."*

# Focus and Solution Classification

We focus on the

profit maximization problem of IaaS cloud providers.

We develop techniques and mechanisms for two main different situations:

When a provider acts solely using their in-house resources to serve customers

When it participates in a cloud federation and benefits from outsourcing requests

# Background (Pricing)

## Pricing definition

The process of determining the rate or fee the provider will receive in exchange for offering services or selling resources

## Pricing Strategies

Marginal-cost pricing, Market-oriented pricing, etc.

## Pricing Factors

Cost of Service

Market Competition

Value to the Customers

## Pricing Models

Usage-based (Pay-as-you-go or consumption-based)

Subscription-based (reservation contract or prepaid scheme)

Demand-oriented

# Background

## Non-storable or perishable commodity

Cloud services similar to many other services, are perishable in nature, and cannot be stored for future

## IaaS providers offer computational services

CPU cycles, network bandwidth, and memory space

Virtual Machine (VM) instances

VMs are non-storable (perishable)

if not utilized at a specific point in time

no value and waste associated underlying data center capacity

## To maximize revenue, IaaS provider adopts differentiated pricing plan with different Quality of Service (QoS)

Usage-based, subscription-based and demand-oriented

# Background (Common Pricing Models)

## On-demand pay-as-you-go instances (usage-based)

- Customers pay for instances based on actual usage

- No long-term commitments

- Billing cycle (e.g., hourly)

- Fixed-rate basis

## Reserved instances (subscription-based)

- Also known as reservation plan

- Reservation fee (premium)

- After which the usage is either free (e.g., GoGrid) or heavily discounted (e.g., Amazon EC2)

- Guaranteed availability.

## Spot instances (demand-oriented)

- Also known as spot market

- Customers submit their bids

- Providers report a market-wide clearing price

- Terminated by the provider if spot market's clearing price rises above the customer's bid

# Background (Cloud Federation)

Resources available in a single data center are limited

A large demand may put pressure in the data center capacity

One possible source for additional resources is idling resources from other providers

Cloud Federation is a collection of individual Cloud providers, which collaborate by trading resources (e.g. computing, storage)

# Outline

Background

Motivations, Research Problem and Thesis Focus

Background

Pricing and Dynamic Pricing

Cloud Federation

## Solutions

Federation of Clouds

Resource Provisioning Policies to Increase Profit

Financial Option Market Model

Single Cloud Provider

Revenue Management with Optimal Capacity Control

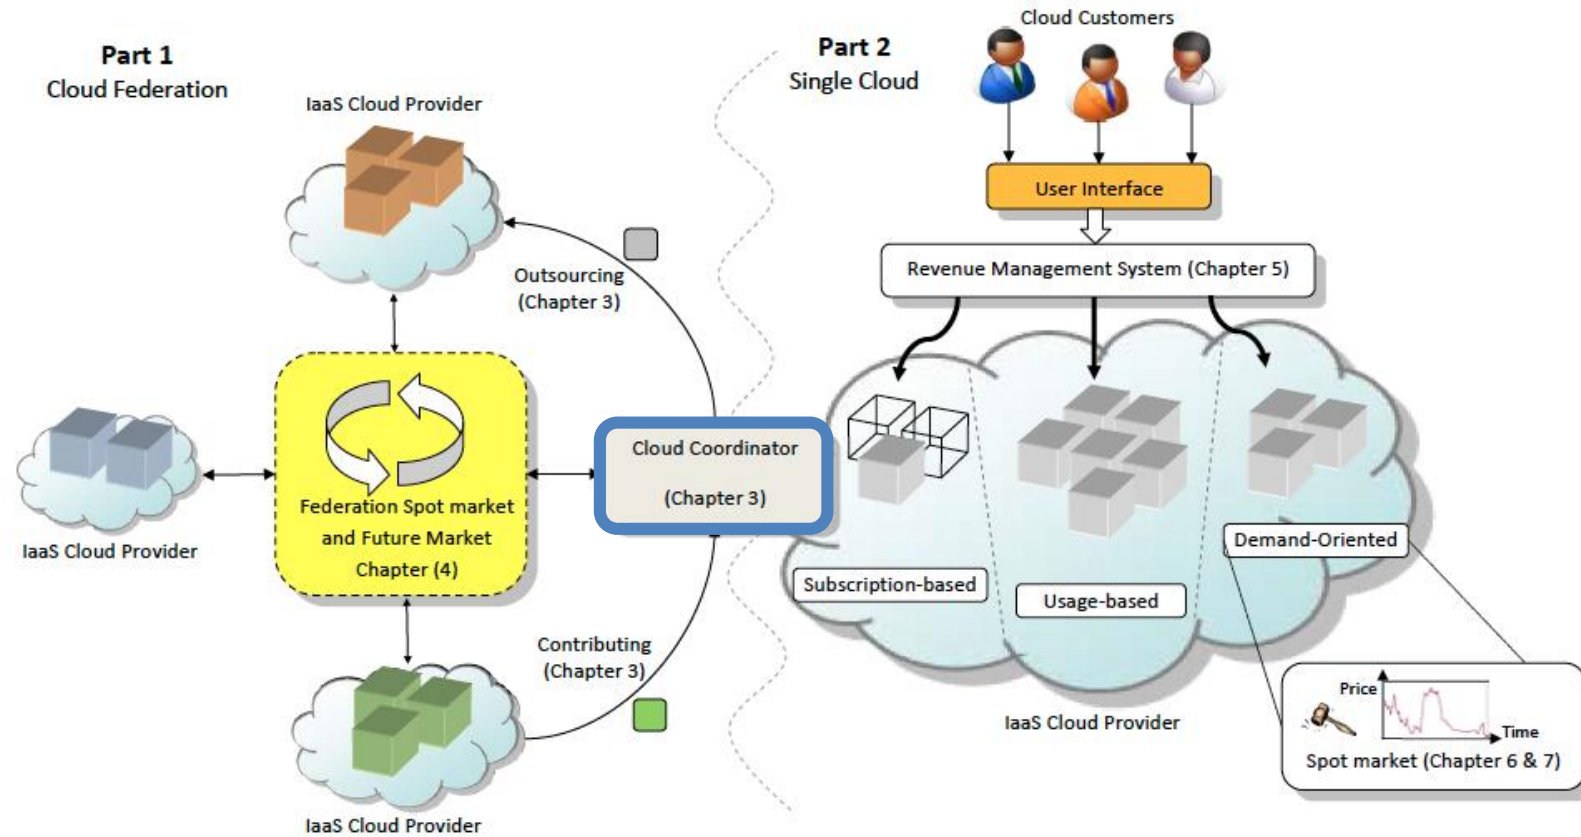
An Auction Mechanism for a Cloud Spot Market

Spot instance pricing as a Service

Conclusions and Future Directions



# Solutions



# Resource Provisioning Policies to Increase Profit

Derived from publication:

Adel Nadjaran Toosi, Rodrigo N. Calheiros, Ruppia K. Thulasiram, and Rajkumar Buyya, "**R**  
**provisioning policies to increase IaaS provider's profit in a federated cloud environm**  
Proceedings of the 13th IEEE International Conference on *High Performance Comput*  
*Communications* (HPCC'11), Banff, Canada, Sep. 2011, pp. 279-287.

# Resource Provisioning Policies in Cloud

Current resource management strategies hinder providers market potential by *limiting the amount of resources* allocated to requests

- QoS is met in a conservative way

- Over-provisioning of datacenter capacity

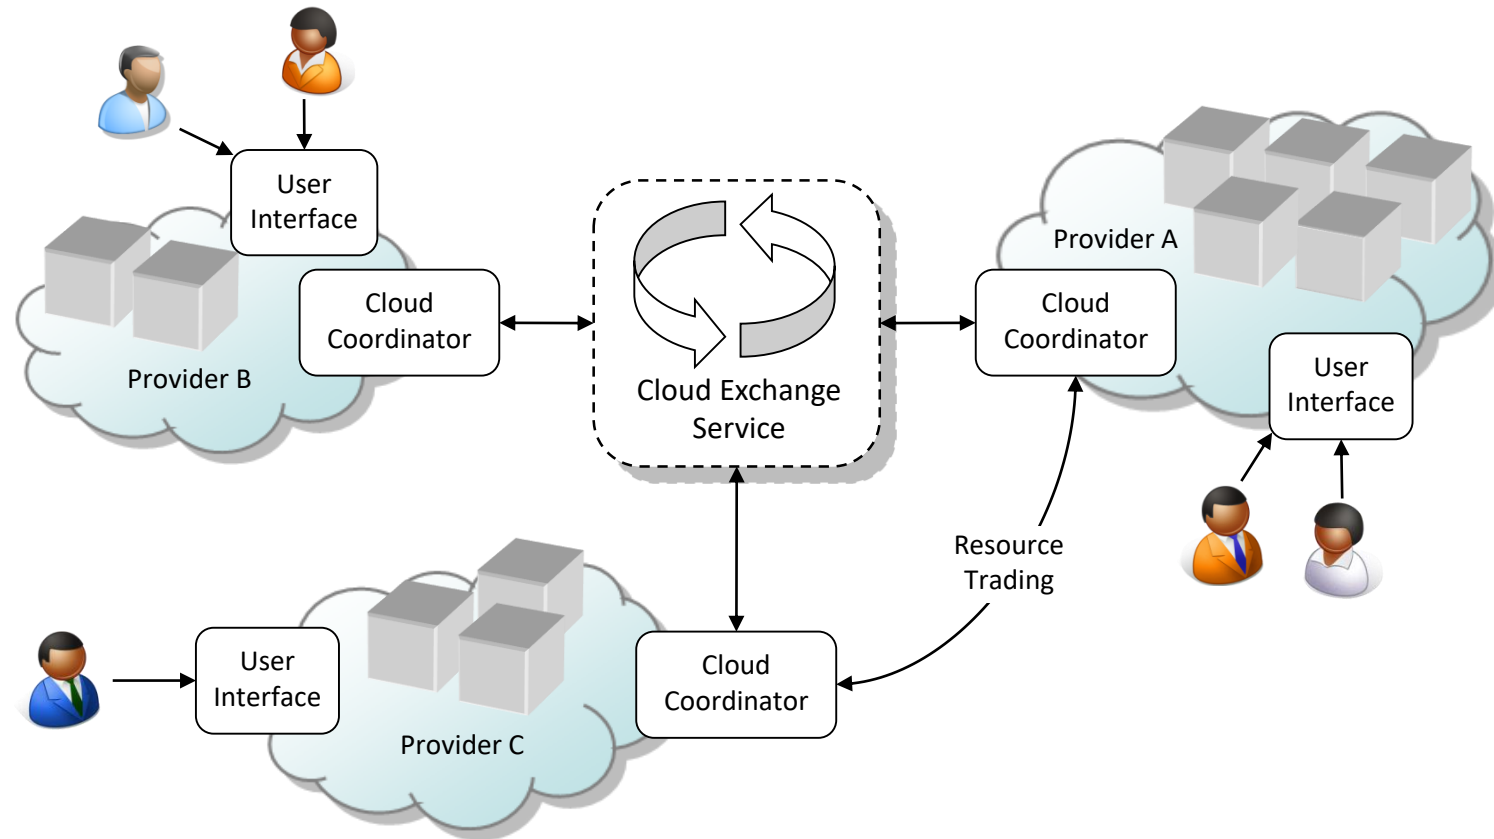
  - average demand of the system is several times smaller than the peak demand

By exploiting cloud federation potentials, providers are able to *dynamically increase the available resources* to serve requests  
*Outsourcing* requests or *renting part of idling resources* to other providers is not a trivial issue

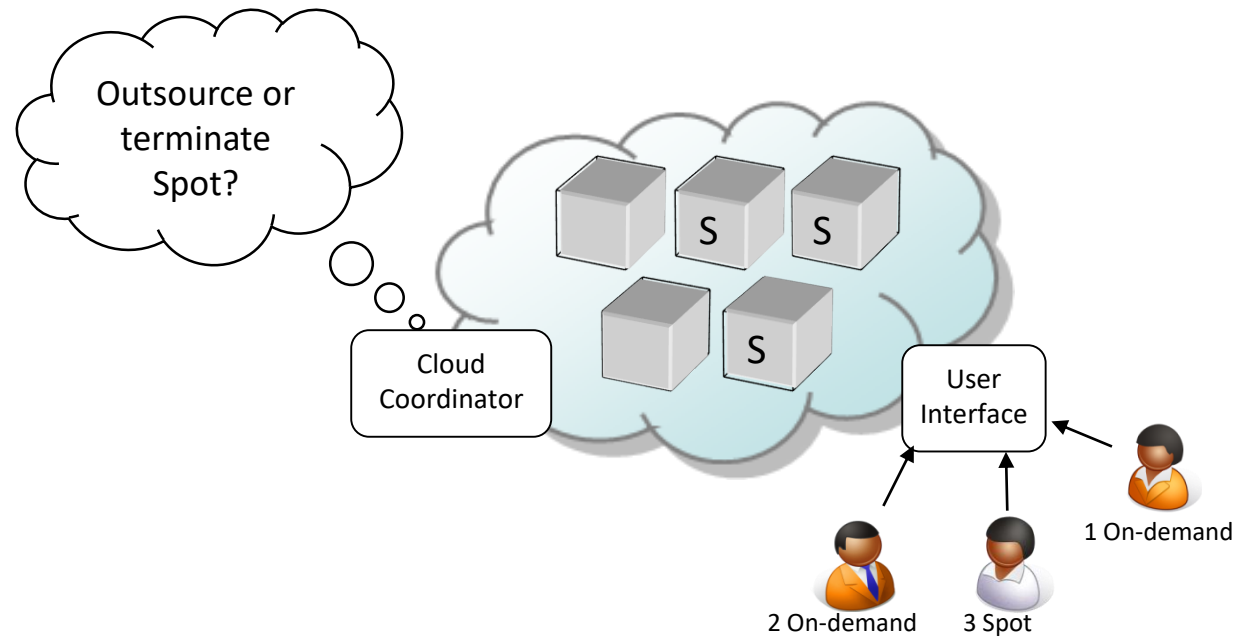
# Research Objectives

Study the *impact of federation* as a mechanism for maximizing cloud provider's profit, utilization and reputation  
Leverage federation potentials by creating *resource provisioning mechanisms* for Cloud providers  
Investigate the effect of applying different *resource trading paradigms* to facilitate federation of providers

# System Model



# Outsourcing decision scenario



Resource provisioning policies for federated-aware providers in the presence of terminable local VMs

# Policies

## Non Federated Totally In-house (NFTI)

Termination of spot VMs with lowest bid

If action does not release enough resources for the new on demand request the request will be rejected

## Federation-Aware Outsourcing Oriented (FAOO)

Fully utilized provider firstly checks the Cloud exchange service for available resources by other members

It outsources the request to the provider that offers the cheapest price

## Federation-Aware Profit Oriented (FAPO)

Based on analytical analysis of instant profit, it decides between outsourcing and termination of spot VMs

# Performance Evaluation - Setup

## Simulation study with CloudSim

CloudSim: A discrete simulator for modelling and simulation of Cloud Computing

## The VM configuration is inspired by Amazon EC2 instances

One VM type (Small Instances: 1 CPU core, 1.7 GB RAM, 1 EC2 Compute Unit, and 160 GB of local storage)

Each datacenter 128 servers, and each one supports 8 VMs.

Due to lack of publicly available workload models and real traces of IaaS Clouds,  
Lublin workload model (one week long simulation)

Each experiment is carried out 20 times

Average of the results is reported.


## Bidding Algorithm:

A uniformly-distributed random value between the minimum of bid \$0.020 and maximum of \$0.085 (on-demand price)


The minimum price is set in such a way that the value offered by customers is still enough to cover operational costs of serving the request



# Results




Impact of load on (a) Profit (b) Utilization (c) Number of rejected on-demand VMs, for a provider with different policies.




Impact of number of providers on (a) Profit (b) Utilization (c) Number of rejected on-demand VMs for a provider with different policies.

# Results(cont.)

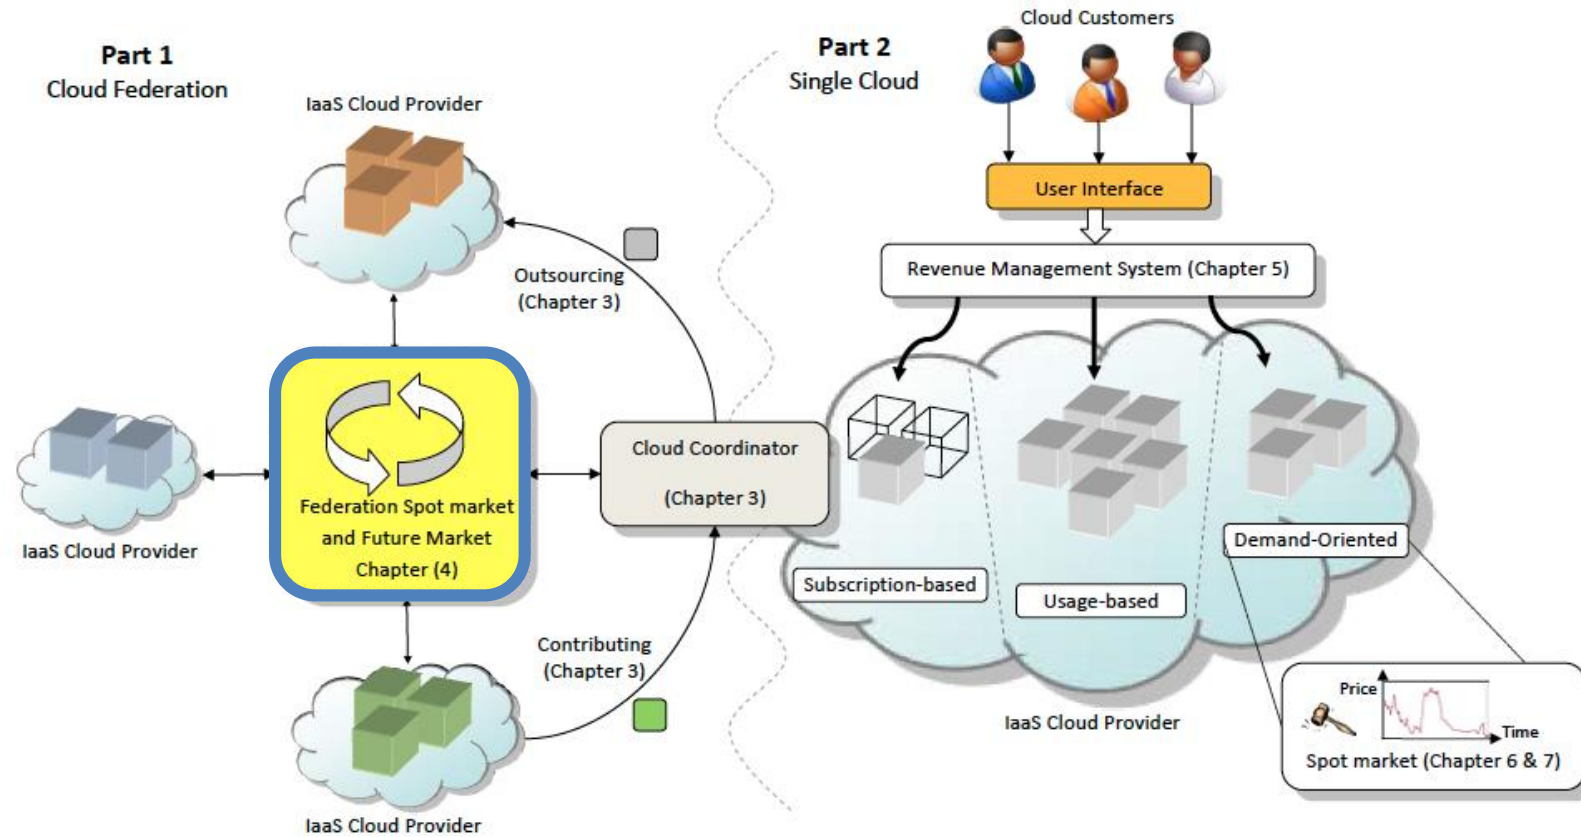


Impact of percentage of spot requests on (a) Profit (b) Utilization (c) Number of rejected on-demand VMs, for a provider with different policies.



Impact of percentage of persistent spot requests on (a) Profit (b) Utilization (c) Number of rejected on-demand VMs for a provider with different policies.

# Solutions



# Financial Option Market Model

Derived from publication:

Adel Nadjaran Toosi, Ruppa K. Thulasiram, and Rajkumar Buyya, "**Financial option market model for federated cloud environments**," in Proceedings of the 5th IEEE/ACM International Conference on *Utility and Cloud Computing (UCC'12)*, Chicago, Illinois, USA, Nov. 2012, pp. 3-12.

# Reserved capacity is not fully utilized!

Huge and unpredictable variation in the load over time for Cloud applications, e.g. web applications.

Economic advantage of using Reserved Instances in comparison with On-Demand Instances, even if they are not fully utilized

Amazon EC2: Using Heavy Utilization Reserved Instances, you can **save up** to 54% for a 1-year term and 71% for a 3-year term vs. running On-Demand Instances.

Reserved Instance Marketplace

All together :

The pattern of utilization at user side causes reserved instances not deployed at all times

# Using Reserved Capacity for On-demand Requests

We are interested to explore the opportunity:

Additional cash flow by releasing underutilized capacity of the reserved instances to the on-demand requests.  
If the unreserved part of the data center experiences high utilization, providers are able to accommodate demand requests on the underutilized reserved capacity of the data center.

Risk of SLA violation!  
Cloud cooperation is a possible solution

# Risks in Cloud Federation Market

## Two more risks:

- a risk of being unable to acquire required resources in the market.

- Short selling resources without having a good knowledge of usage loads and hence violating the C

- Future price variations in the federation market using past price history

A financial option-based market model is introduced for  
federation of Cloud providers

which helps providers increase their profit and mitigate the risks (risks of violating QoS or paying extra

# Financial Option Contract

A financial option is a contract for a future.

Transaction between two parties:

Holder and seller of the contract.

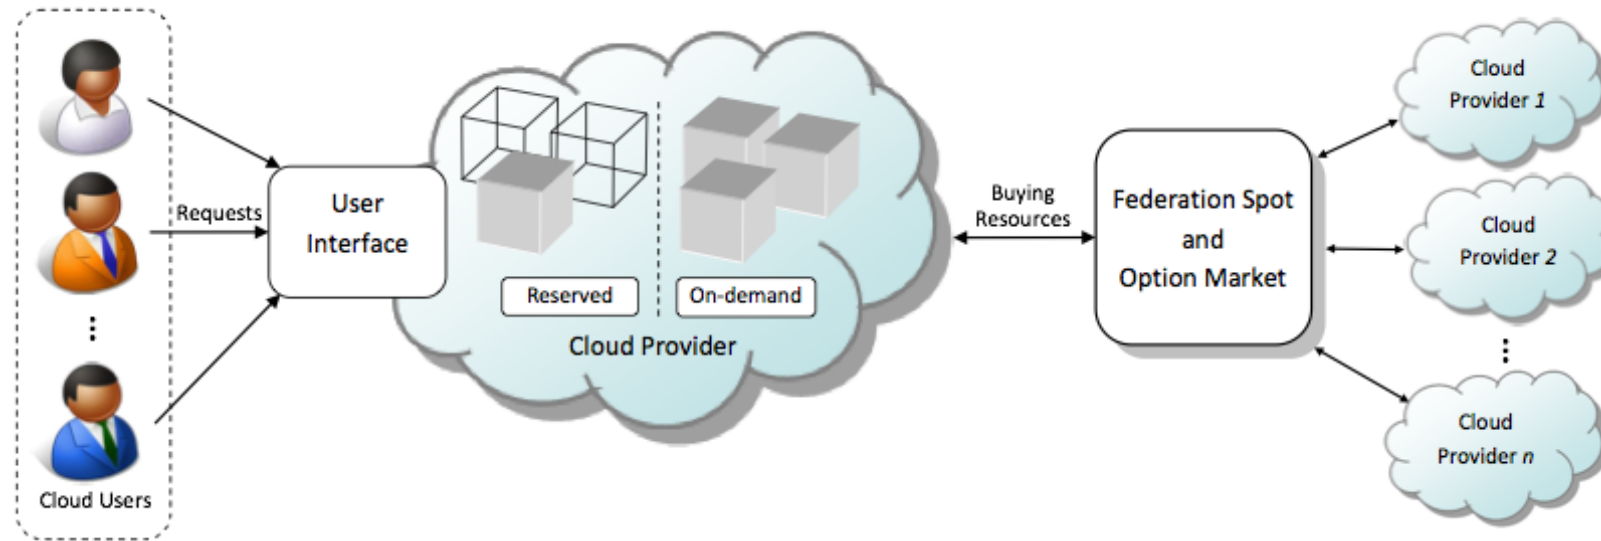
A financial option gives the holder the right, *but not the obligation*, to (or to sell) an underlying asset at a certain price, called the *strike price* (exercise price), within a certain period of time, called *maturity date* (expiration date).

The seller is obligated to fulfill the transaction.


As a compensation the seller collects an upfront payment at the beginning of the contract, called *premium*.



# The System Model



# When to buy option?



# Policies

## Baseline In-house Isolated Pool Policy (IIP)

Provider works independently, without participating in the federation.

Isolated pools of physical nodes for on-demand and reserved instances .

## Baseline Federated Isolated Pool Policy (FIP)

If the provider is not able to serve on-demand requests locally it outsource them through federation spot market.

To be always cost-effective, if the spot price in the federation market is higher than the local on-demand price then the provider rejects on-demand request.

The pool of physical nodes for reserved and on-demand instances are isolated to prevent rejection of reserved requests.

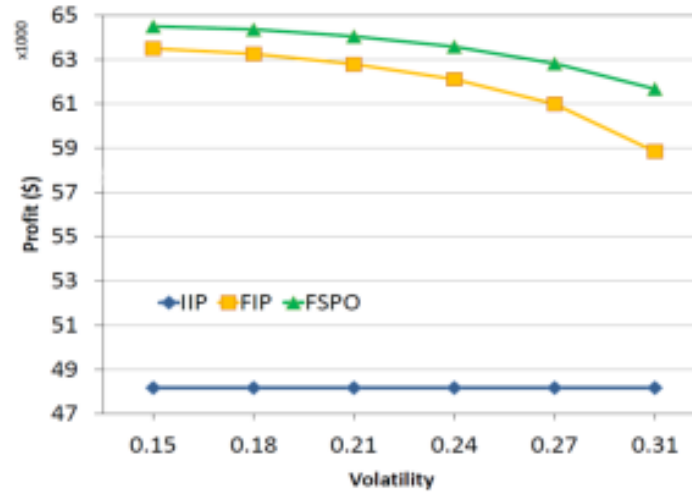

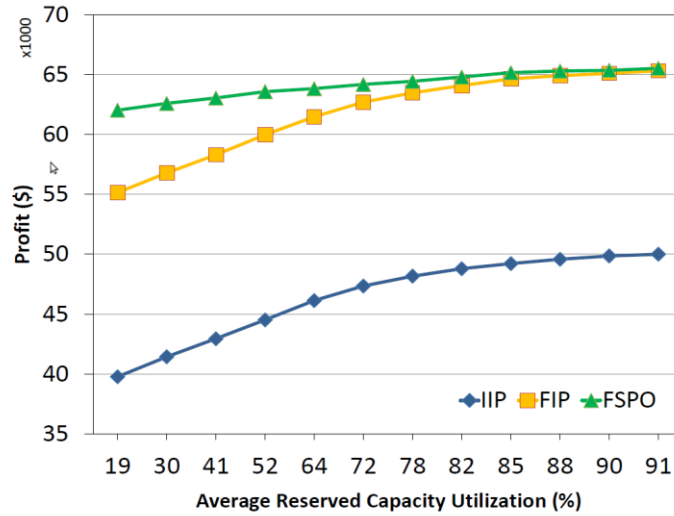
## Federated Shared Pool Option-Enabled Policy (FSPO)

The provider accommodates excess on-demand requests in the underutilized reserved capacity.

The provider buys an option whenever he accommodates on-demand requests in the reserved capacity.

If a reserved request comes in and the provider is not able to serve it locally, the option is exercised and the reserved request is out of the strike price of the option.

# Experimental Results



# Number of Request Rejections

NUMBER OF ON-DEMAND (O), RESERVED (R), REJECTED  
ON-DEMAND (RO), REJECTED RESERVED (RR), OUTSOURCED  
ON-DEMAND (OO), AND OUTSOURCED RESERVED(OR) REQUESTS  
FOR THE PROVIDER WITH POLICIES.

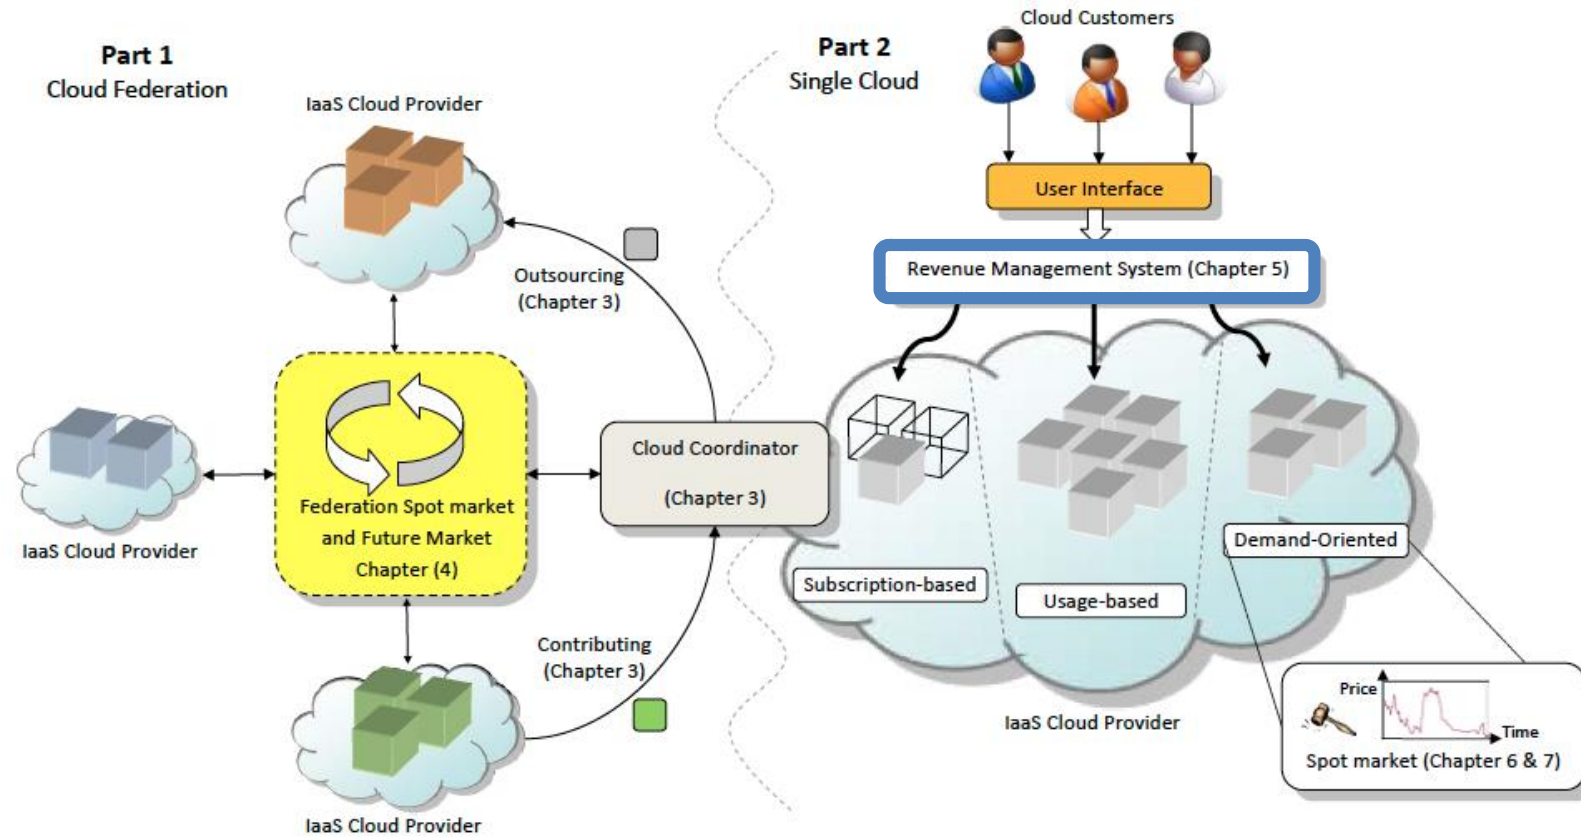
Policy	O	R	RO	RR	OO	OR
IIP	106340	36590	42450	0	0	0
In-house+ Shared Pool	106340	36590	19188	6636	0	0
FIP	106340	36590	219	0	42230	0
Federated+ Shared Pool	106340	36590	111	38	19078	6598
FSPO	106340	36590	111	0	19078	6636

# Effect of Maturity Time

IMPACT OF OPTION MATURITY TIME ON PROVIDER'S PROFIT USING FSPO.

Maturity time (Day)	Bought Option	Profit (\$)
7	3086	63588
10	2612	63588
30	1913	63583
60	1745	63578
90	1701	63574

# Solutions



# Revenue Management with Optimal Capacity Control

Derived from:

Adel Nadjaran Toosi, Kurt Vanmechelen, Ramamohanarao Kotagiri, and Rajkumar Buyya, "Revenue  
**Maximization with Optimal Capacity Control in Infrastructure as a Service Cloud M**  
*IEEE Transactions on Cloud Computing* (TCC), 2014, in



# IaaS providers' various pricing plans (or markets)

## on-demand pay-as-you-go

- Generates highest revenue per unit capacity

- Demand uncertainty

## Reservation (subscription)

- Risk-free income from reservations

- Guaranteed cash flow through long-term commitments

- Compensate for the demand uncertainty of on-demand pay-as-you-go

- Generates lower revenue per unit capacity

- The provider is liable to offer guaranteed availability to honor SLA

## Spot market

- Selling spare capacity in the data center

- Attract price-sensitive users that are capable of tolerating service interruptions

- Without being exposed to the risks resulting from overbooking capacity

# Problem definition

The use of multiple pricing plans introduces:

Non-trivial trade-offs in revenue maximization

Our main **research question** is the following:

"with **limited resources** available, and considering the **dynamic** and **stochastic** nature of customers' **demand**, how can **expected revenue** be maximized through an **optimal allocation** of capacity to each pricing plan?"

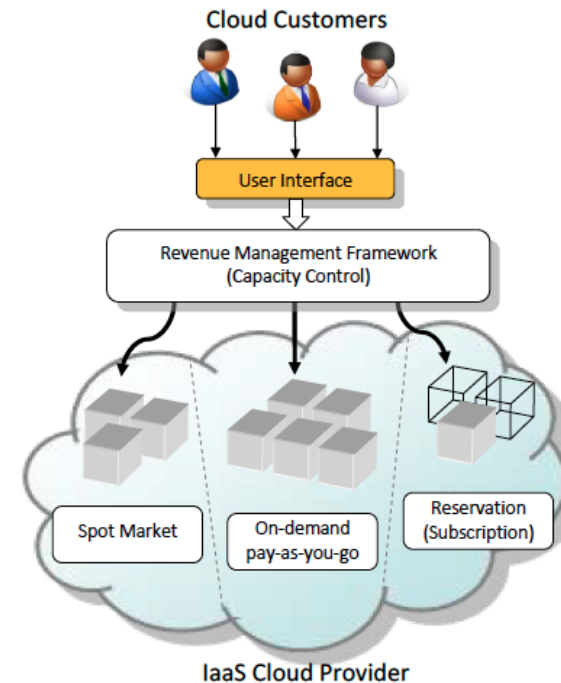
**Assumptions:**

Customers do not fully utilize their reserved capacity in the reservation's lifetime

We do not allow accommodating on-demand requests in underutilized reserved capacity

Creates the risk of SLA violations.

Provider accommodates spot instances in the underutilized reserved capacity of the data center



# Our Solution

## We frame our algorithmic contributions

- Within a revenue management framework

- Supports the three presented pricing plans

- Incorporates an admission control system for requests of the reservation plan.

## We formulate the optimal capacity control problem

- As a finite horizon Markov Decision Process (MDP)

- A stochastic dynamic programming technique

  - To compute the maximum number of reservation contracts the provider must accept

  - For a large capacity provider the stochastic dynamic programming technique is computationally prohibitive.

- We therefore present two algorithms to increase the scalability of our solution

# Contributions

## Proposed Algorithms

### Optimal Algorithm

- A stochastic dynamic programming technique

- As a finite horizon Markov Decision Process (MDP).

### Pseudo Optimal Algorithm

- Based on optimal algorithm

- Only increases the spatial and temporal granularity of the problem

- To solve it in a time suitable for practical online decision making

### Heuristic Algorithm

- Sacrifices accuracy to an acceptable extent to increase scalability

- Through a number of simplifying assumptions on

  - Reserved capacity utilization and

  - The lifetime of on-demand requests.

# Pseudo Optimal

---

## Algorithm 1 Pseudo Optimal Algorithm

---

**Input:**  $t, l_t^r, l_t^o, i_t$

**Output:**  $maxrev$

```

1:  $dp \leftarrow \{-1\}$  ▷ matrix  $dp$  is used for memoization and all cells are initialized with -1.
2: function  $V(t, l_t^r, l_t^o, i_t)$ 
3:   if  $dp[t][l_t^r][l_t^o][i_t] \neq -1$  then
4:     return  $dp[t][l_t^r][l_t^o][i_t]$ 
5:   end if
6:   if  $t = \tau$  then
7:      $dp[t][l_t^r][l_t^o][i_t] = 0$ 
8:     return 0
9:   end if
10:   $maxrev \leftarrow 0$ 
11:  for  $r_t \leftarrow 0$  to  $\min(C - l_t^r - l_t^o, d_t^r)$  do
12:     $rev \leftarrow 0$ 
13:     $l_{t+1}^r \leftarrow l_t^r + r_t - e_t^r$ 
14:     $o_t \leftarrow \min(C - l_t^r - l_t^o - r_t, d_t^o)$ 
15:     $s_t \leftarrow \min(C - (l_t^r + r_t)u_{i_t} - l_t^o - o_t, d_t^s)$ 
16:     $\lambda \leftarrow (\tau - t) / \tau$ 
17:     $\gamma(\zeta_t, r_t) \leftarrow B\lambda r_t \varphi + BT(\alpha p(l_t^r + r_t)u_{i_t} +$ 
       $p(l_t^o + o_t) + \beta p s_t)$ 
18:    for  $l_{t+1}^o \leftarrow 0$  to  $l_t^o + o_t$  do
19:      for  $i_{t+1} \leftarrow 0$  to  $|U|$  do
20:         $P(\zeta_{t+1}|\zeta_t, r_t) \leftarrow \text{Bin}(l_{t+1}^o, l_t^o + o_t, q) \times P(u_t = u_{i_{t+1}})$ 
21:         $rev \leftarrow rev + \gamma(\zeta_t, r_t) + P(\zeta_{t+1}|\zeta_t, r_t) \times V(t+1, l_{t+1}^r, l_{t+1}^o, i_{t+1})$ 
22:      end for
23:    end for
24:    if  $rev \geq maxrev$  then
25:       $maxrev \leftarrow rev$ 
26:    end if
27:  end for
28:   $dp[t][l_t^r][l_t^o][i_t] \leftarrow maxrev$ 
29:  return  $maxrev$ 
30: end function

```

---

## Details of *Optimal Algorithm*

Computational Complexity

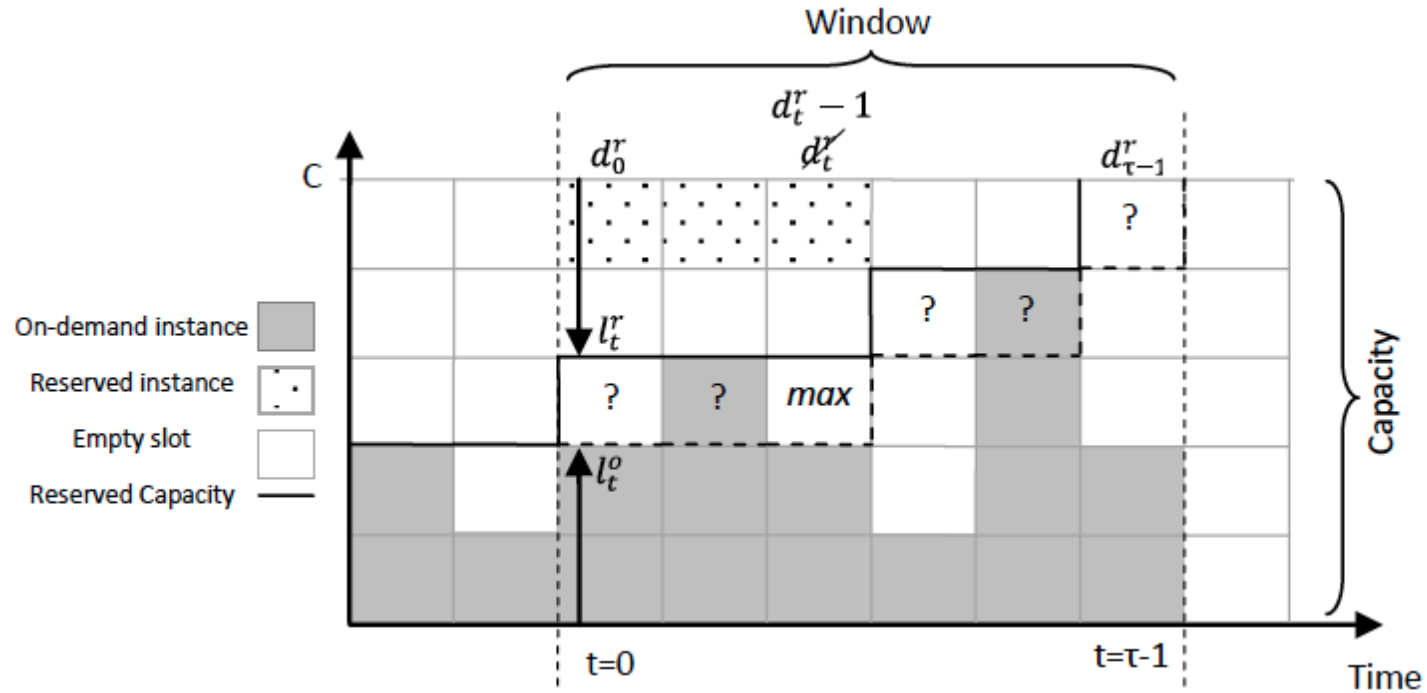
B: The number of instances per block of capacity

e.g., B = 100 VMs

T: The number of billing cycles per time slot

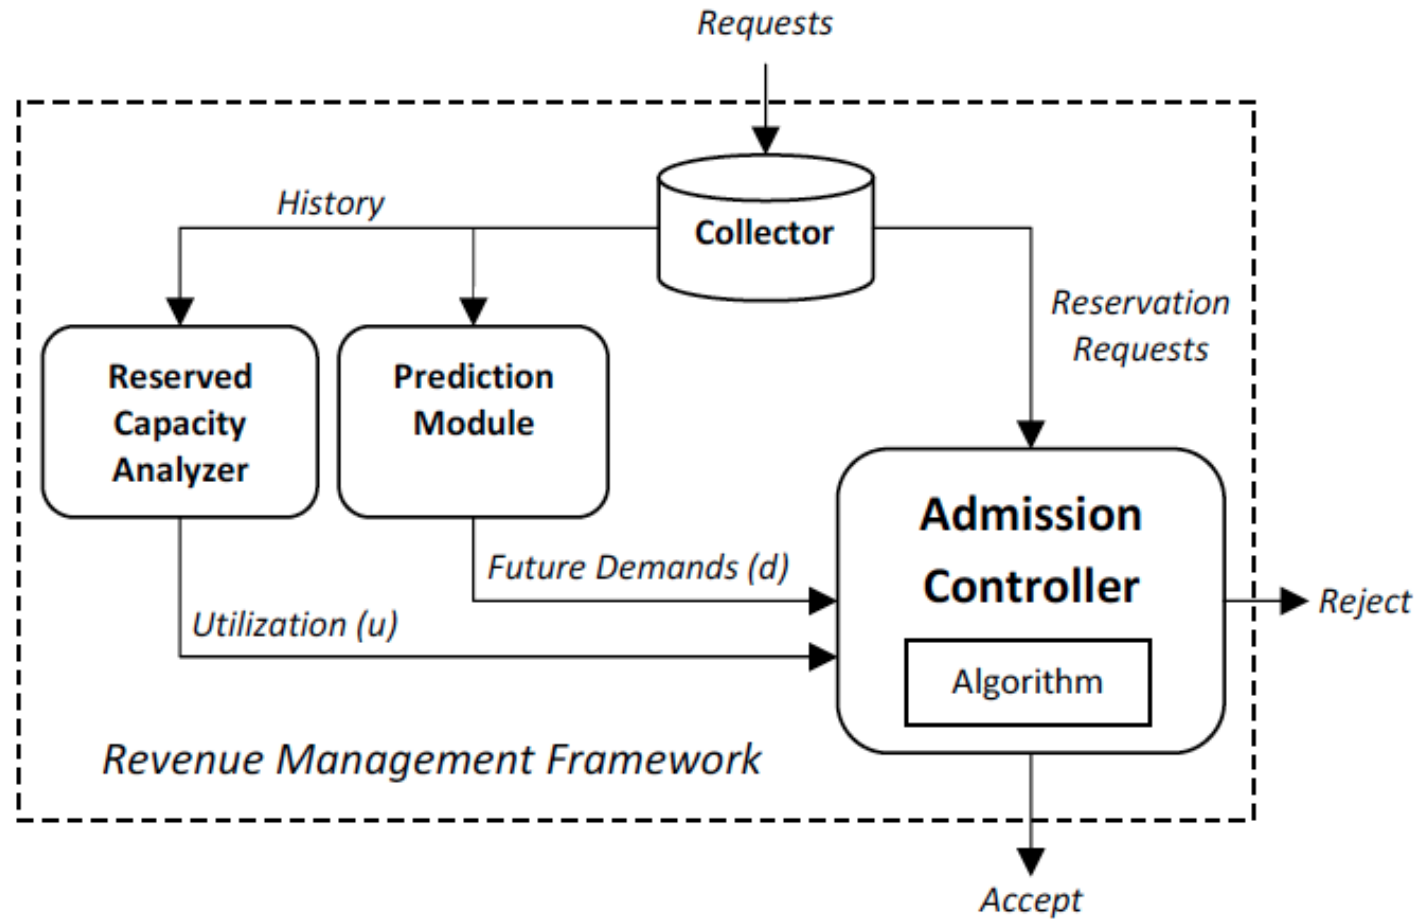
e.g., T = 168 hours.

# The heuristic algorithm



Computational Complexity:  $O(\tau \times C)$

# Key modules of the revenue management framework



# Performance Evaluation

We evaluate our proposed framework through:

- Large-scale simulations (12 months)

- Driven by cluster-usage traces that are provided by Google.

No publicly available workload traces of real-world IaaS clouds:

- We propose a **scheduling algorithm** that generates VM requests based on the user resource usage in these traces.

Pricing conditions that are aligned with those of Amazon EC2

Benchmark Algorithm:


- no admission control referred to as *no-control*

Simulation Environment

- Extended CloudSim (pricing plans and the proposed revenue management framework)

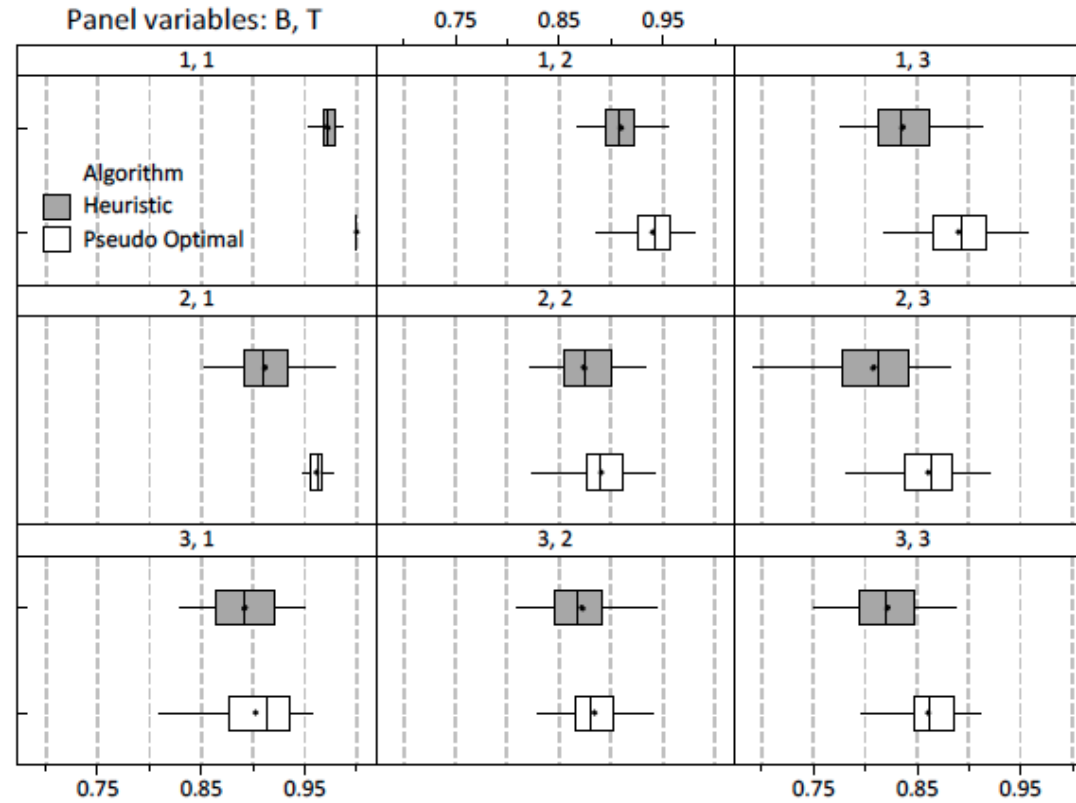


# Experimental Results (1)



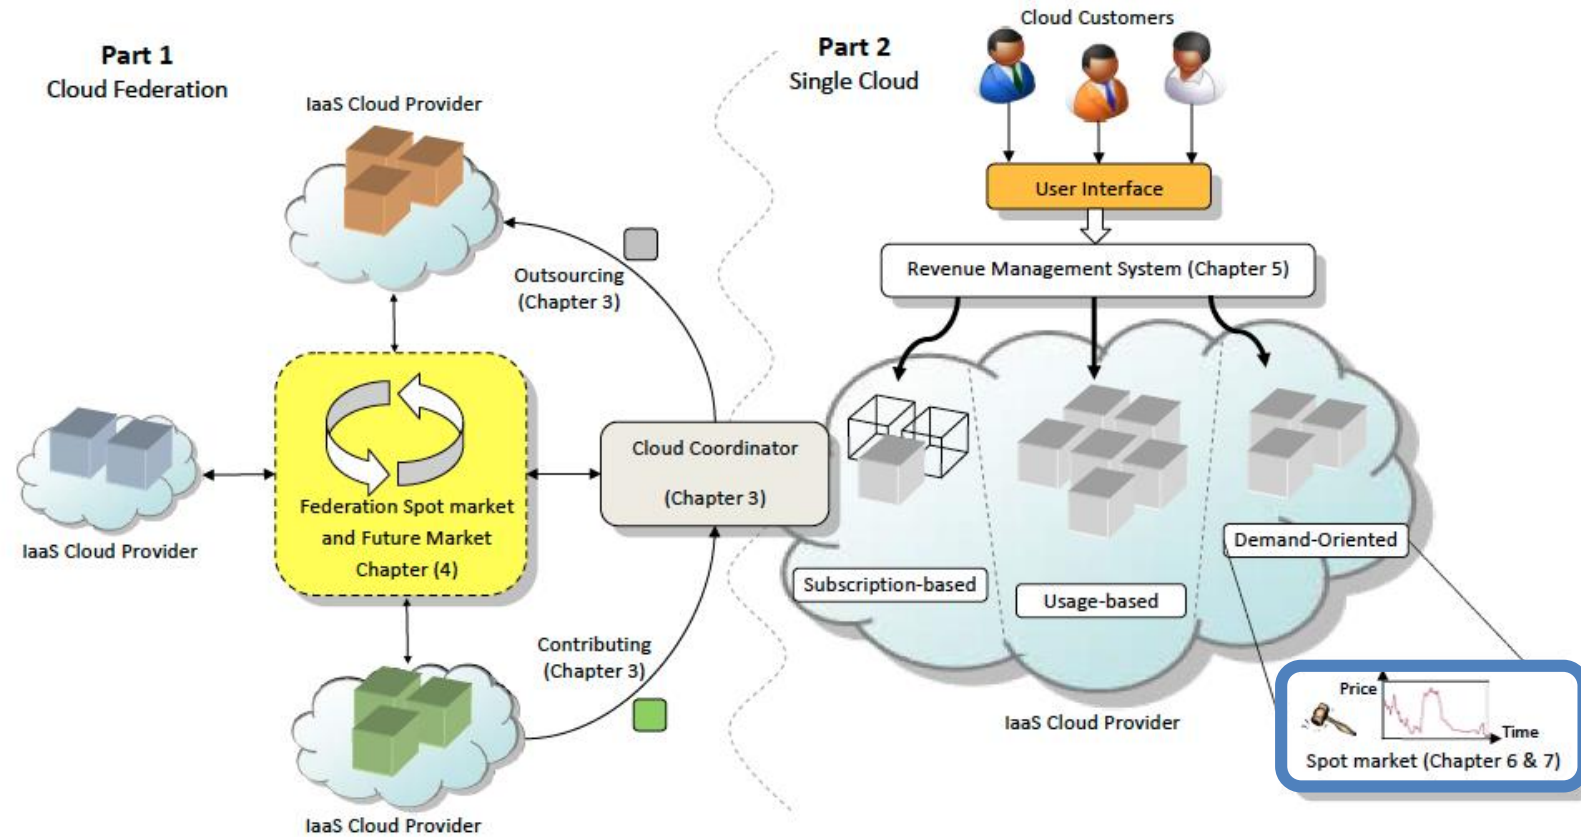
The revenue performance of the proposed revenue management framework under different algorithms normalized to the outcome of no-control algorithm ( $B = 100$  and  $T = 75$ ).

# Experimental Results (2)



The revenue performance of the pseudo optimal and heuristic algorithms with different values of B and T. All values are normalized to the outcome of the optimal solution.

# Solutions



# An Auction Mechanism for Cloud Spot Markets

Derived from:

Adel Nadjaran Toosi, Kurt Vanmechelen, Farzad Khodadadi, and Rajkumar Buyya, "**An Auction Mechanism for Cloud Spot Markets**," *ACM Transactions on Autonomous and Adaptive Systems*, Volume 11, Issue 1, 2016.

# Motivation

Computational resources sold by a cloud provider can be characterized:

As a non-storable or perishable commodity

Demand for computational resources is:

non-uniform over time

These motivates the use of **dynamic forms of pricing** in order to optimize revenue

Well-designed auction mechanisms are particularly effective

# Why Auctions?

Incentivize users to reveal their **true valuation** of the service (i.e., report the price they are willing to pay for resources)

Ensure resources are allocated to those who value them the most

Correctly price resources in line with **supply and demand** conditions by creating competition among buyers.

# Background

## Auction:

A common market mechanism with a set of rules **determining prices** and **resource allocations** on basis of bids submitted from the market participants.

## Auctions can be in assorted shapes and have different characteristics:

Single dimensional (e.g., only bid price) or multi-dimensional (e.g., bid price plus quantity),

Single-sided (e.g., only customers submit bids) or double-sided (e.g., double auction),

Open-cry or sealed-bid,

Single-unit (e.g., a single good or service) or multi-unit (e.g., multiple units of the good),

Single item (e.g., one type of service) or multi-item (e.g., combinatorial auction).

## Goals in designing auction:

Truthfulness, Revenue maximization, Allocative efficiency, Fairness

# Background

## Optimal mechanism design

When goal of mechanism is to maximize the profit or revenue for the seller (provider)

## Optimal mechanism design can be categorized in

### Bayesian optimal mechanism design

the valuations of the participants in the auction are drawn from a known prior distribution

based on the seminal work by Myerson

### Prior free optimal mechanism

Determining the prior distribution is not practical, convenient or even possible in advance

Digital Goods: Competitive Framework by Goldberg and Hartline



# Problem Definition and Goals

## Amazon Web Services (AWS)

Spot Instances (Auction-like mechanism)

Customers communicate their bids a VM instance hour to AWS.

AWS reports a market-wide spot price at which VM instance use is charged,

while terminating any instances that are executing under a bid price that is lower than the market price.

## AWS has revealed no detailed information regarding

Their auction mechanism and

The calculation of the spot price.

## The design of an auction mechanism

Efficient, fair, and profit-maximizing

Open research challenge,

Of great interest to cloud providers.

# Contributions

We design an auction mechanism

aimed at maximizing profit from selling the spare capacity of non-storable resources available in cloud data centers

A multi-unit, online recurrent auction, two-dimensional bid domain

Extension of Consensus Revenue Estimate (CORE) mechanism

Work by Goldberg and Hartline

Envy-free, truthful with high probability and generates near-optimal profit for the provider.

# Contributions (cont.)

## Benchmark Algorithm: An optimal auction mechanism (HTAOPT)

- Dynamic programming

- To quantify the efficiency loss caused by the lack of information on the amount of time a bidder wants to run a VM.

## A method for dynamically computing

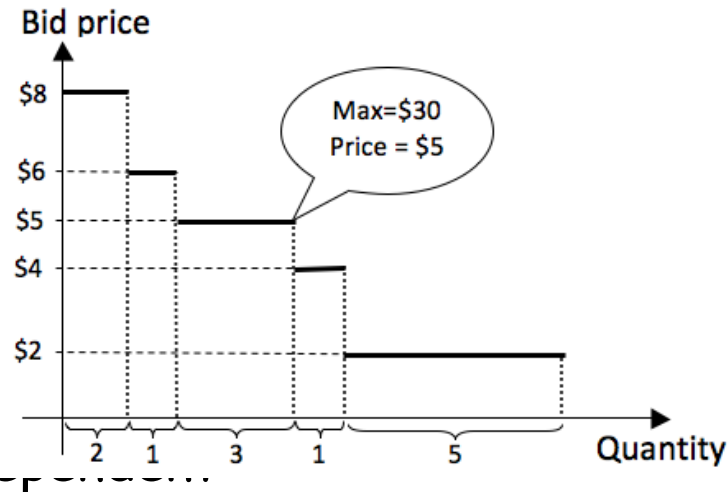
- A reserve price

- The lowest possible price that the provider accepts.

- Based on a coarse grained data center power usage model

# Research Problem

Optimal Single Price Auction:  $\mathcal{F}(\mathbf{d})$



The order-indeper  
Optimal order-ind

Not single price (i.e., not envy-free)

It is not fair

$\mathcal{F}(\mathbf{d}_{-i})$

How to compute a single price for an order-independent auction while attaining the revenue of the optimal auction?

# Consensus Revenue Estimate

We are interested in a mechanism that provides us with a sufficiently accurate estimate of  $\mathcal{F}(\mathbf{d})$

While the estimate is constant on  $\mathbf{d}_{-i}$  for all  $i$  (i.e. achieves consensus).  
 is limited by a constant fraction of  $\mathcal{F}(\mathbf{d})$

it is possible to pick a good estimate

We consider a limit on maximum quantity ( $r$ )

Let  $r$  denote the supremum of the number of requested units in  $\mathbf{d}$

Let  $m$  be the number of sold units in  $F$

$$\frac{m-r}{m} \mathcal{F}(\mathbf{d}) \leq \mathcal{F}(\mathbf{d}_{-i}) \leq \mathcal{F}(\mathbf{d})$$

We use Consensus Revenue Estimate

# Revenue Extraction Mechanism

Cost sharing mechanism proposed by Moulin and Shenkar  
Extracts a target revenue from the set of bidders if this is possible

It finds the  $k$  bidders with the highest bid values that allow for the extraction of  $R$  (target revenue)

Given a target revenue  $R$ , the revenue extraction mechanism is truthful for the price dimension but not for the quantity dimension.

# Reserve Price

The reserve price for most perishable goods and services is considered low at their expiration time

The reserve price for flight seats is theoretically negligible.

A significant part of the service cost in cloud data centers is related to power consumption of physical servers.

We propose a method for calculating **dynamic reserve prices** based on a cost model that incorporates data center **PUE (load, outside temperature)** and **electricity cost**.

# The Online Ex-CORE Auction

---

## Algorithm 3 The Online Ex-CORE Auction

---

**Input:**  $\mathbf{d}, p_{cur}, p_{optprv}$   $\triangleright$   $\mathbf{d}$  is the list of orders, sorted in descending order of bids,  $p_{cur}$  is current market price,  $p_{optprv}$  is the optimal single price in the previous round.

**Output:**  $p$   $\triangleright$  Sale Price

```
1:  $p_{opt} \leftarrow opt(\mathbf{d})$ 
2: if  $p_{opt} = p_{optprv}$  then
3:   return  $p_{cur}$ 
4: end if
5:  $r \leftarrow$  the largest  $r_i$  in  $\mathbf{d}$ 
6:  $m \leftarrow \underset{\sigma_i(\mathbf{d})}{\operatorname{argmax}} b_i \sigma_i(\mathbf{d})$ 
7: if  $m \leq r$  then
8:   return  $p_{opt}$   $\triangleright$  single optimal price
9: else
10:   $\rho \leftarrow \frac{m}{m-r}$ 
11:  Find  $c$  in  $\rho \ln(c) + \rho - c = 0$ 
12:   $u \leftarrow \operatorname{rnd}(0, 1)$   $\triangleright$  chosen uniformly random on  $[0, 1]$ 
13:   $l \leftarrow \lfloor \log_c(\mathcal{F}(\mathbf{d})) - u \rfloor$ 
14:   $R \leftarrow c^{(l+u)}$ 
15:   $j \leftarrow$  the largest  $k$  such that  $\frac{R}{\sigma_k(\mathbf{d})} \geq b_k$ 
16:  return  $\frac{R}{\sigma_j(\mathbf{d})}$ 
17: end if
```

---



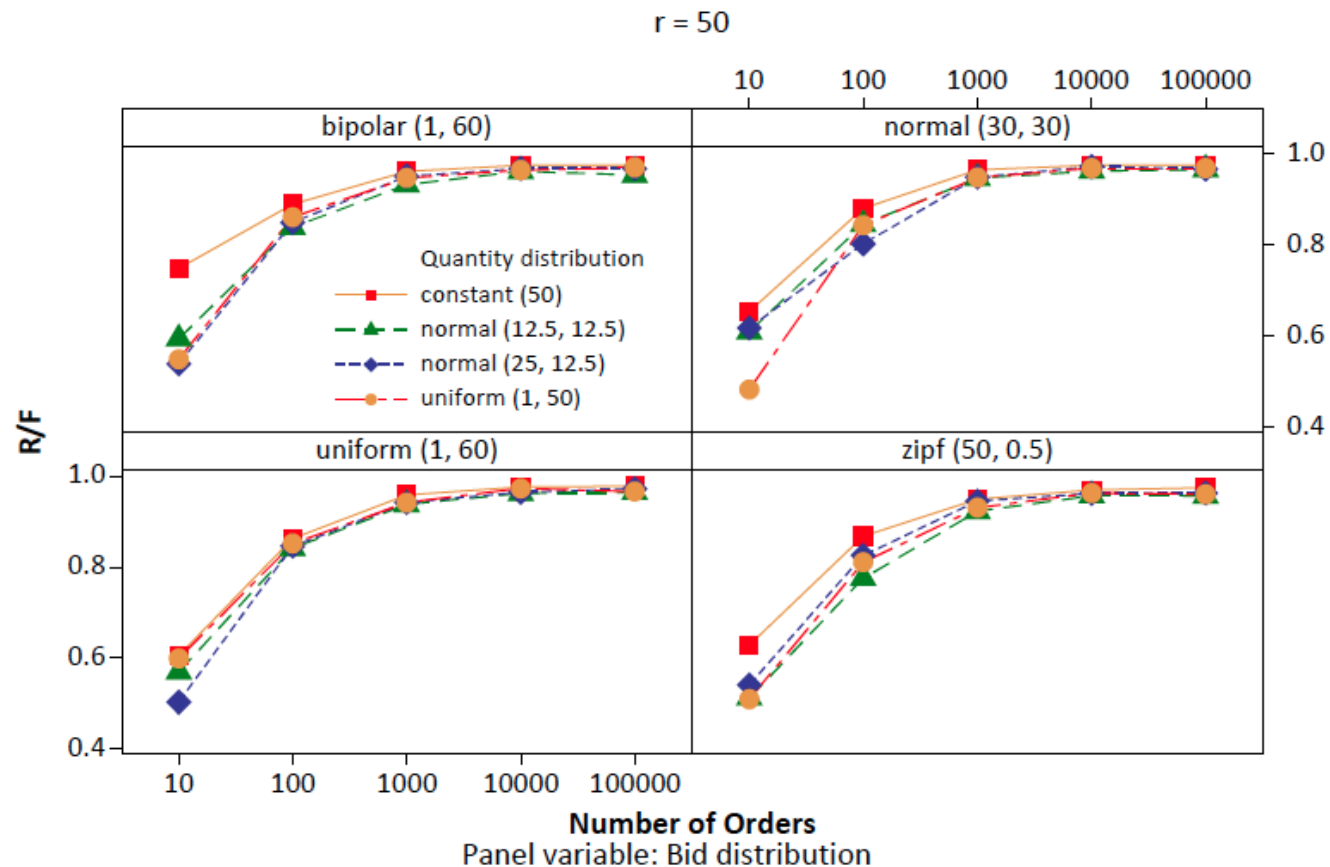
# Benchmark Algorithms

Optimal Single Price Auction (OPT)

Holding Time Aware Optimal Auction (HTA-OPT)

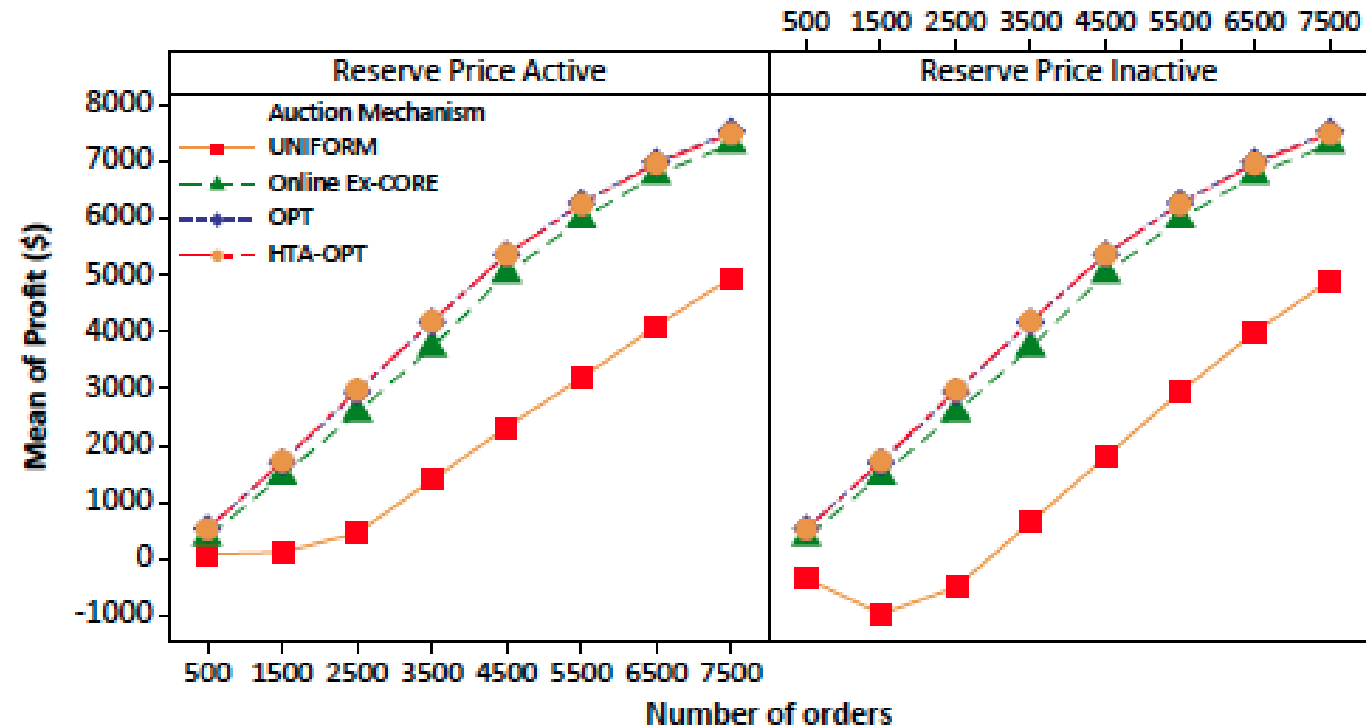
Uniform Price Auction

# Performance Evaluation (Single Round)




Ratio of gained revenue by the Ex-CORE auction to optimal auction under different distribution of orders

# Performance Evaluation (Online)



Average profit gained and with different auction mechanisms.

# Results



Reserve price (green dashed line) and spot market price generated by online Ex-CORE (blue solid line) in a sample simulation run when the number of orders is 1500.

# SipaaS: Spot instance pricing as a Service

Derived from:

Adel Nadjaran Toosi, Farzad Khodadadi, and Rajkumar Buyya, "**SipaaS: Spot instance pricing as a Service and its Implementation in OpenStack**," Software Software: Practice and Experience, Wiley Press, 2014 (in preparation).

# Spot instance pricing as a Service

## The implementation of the proposed auction mechanism

by identifying a framework called Spot instance pricing as a Service (SipaaS).

### SipaaS:

is an open source project offering a set of web services to price and sell VM instances in a spot market.

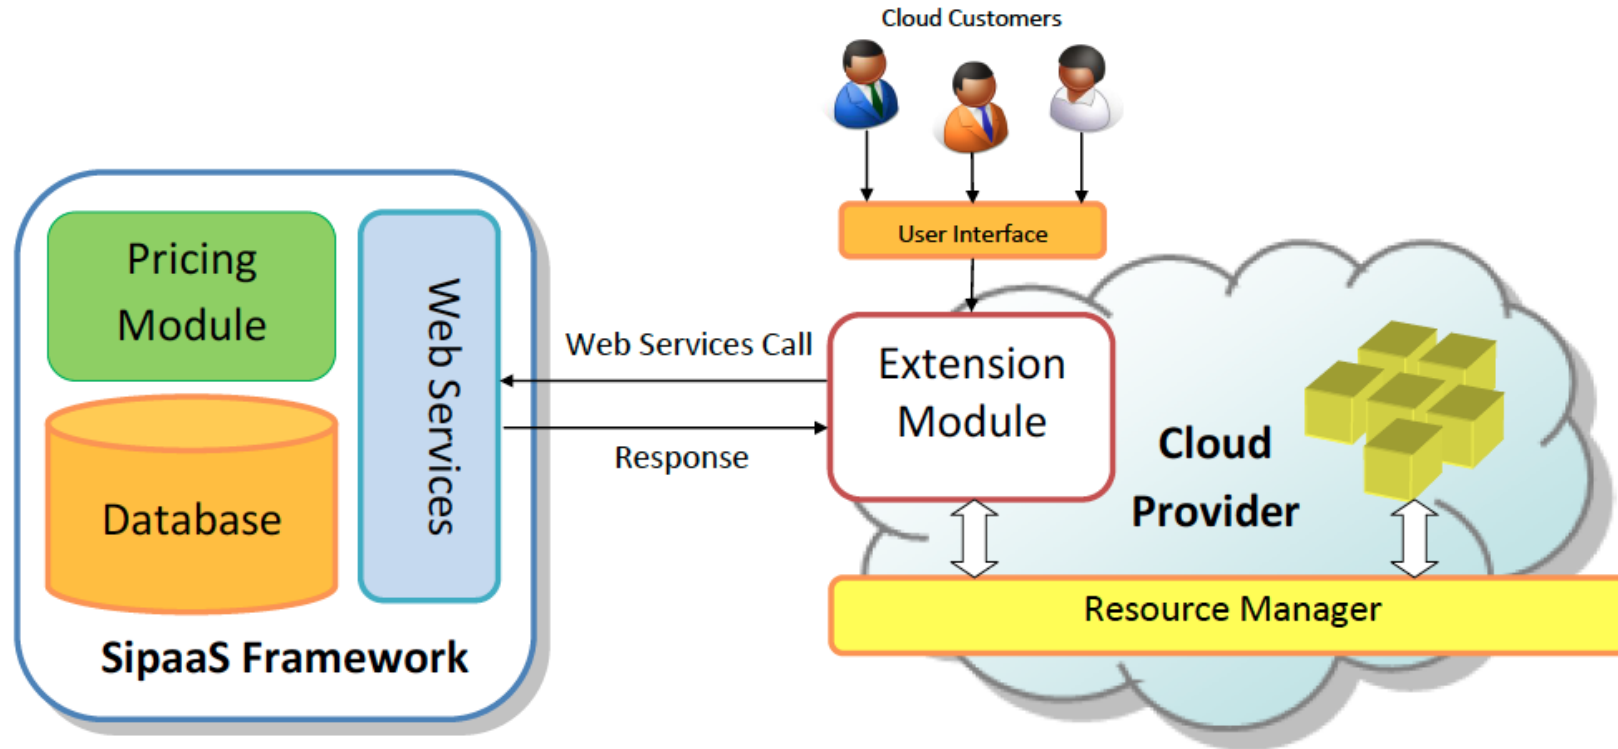
### Utilizing SipaaS

Cloud providers require installation of add-ons in their existing platform to make use of the framework.

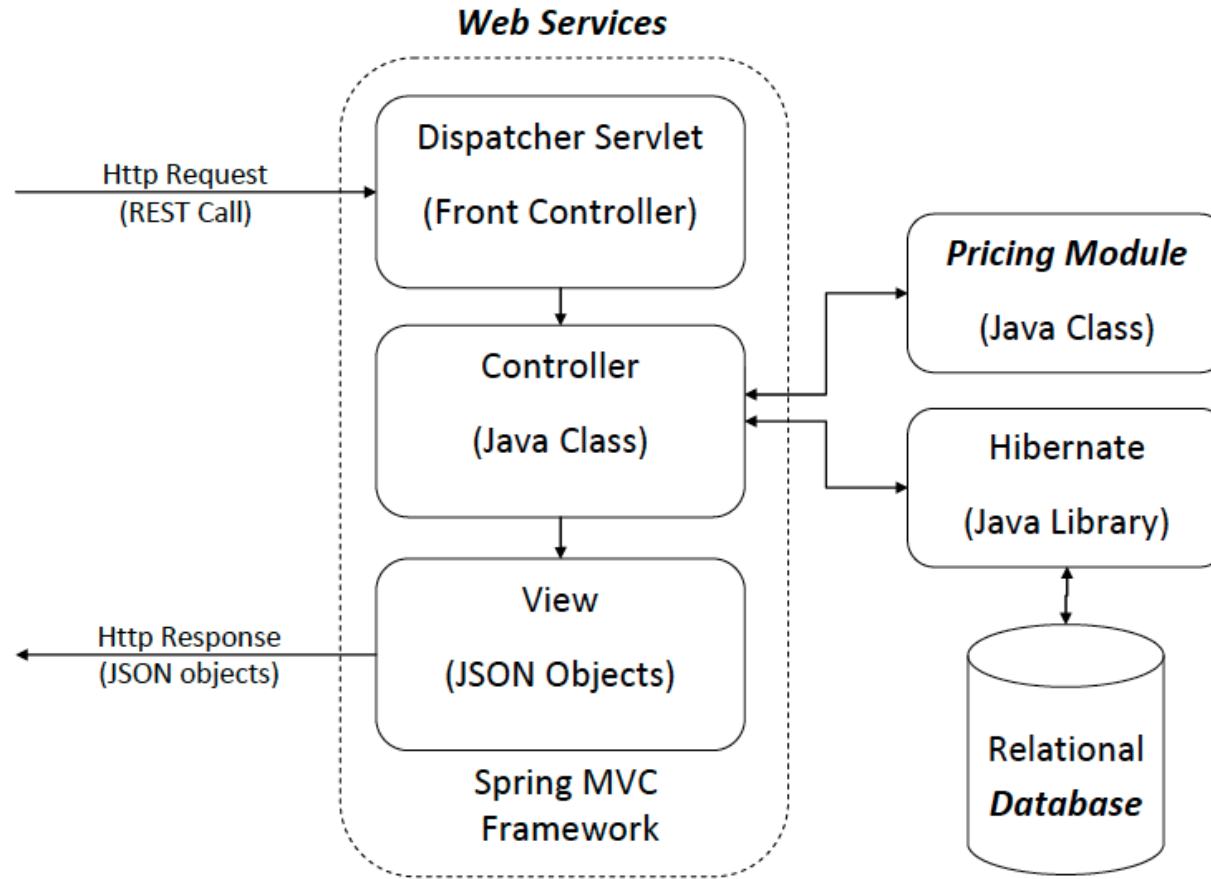
An extension to the Horizon –the OpenStack dashboard project

To employ SipaaS web services and to add a spot market environment to OpenStack.

# System Model



# SipaaS Architecture





# List web services

Web Service Name	Input parameter(s)	Output
register	name	credentials
unregister	accesskey, secretkey	status
regvmttype	accesskey, secretkey, type	status
unregvmttype	accesskey, secretkey, type	status
setavailables	accesskey, secretkey, vmttype, quantity	price
setmaxqty	accesskey, secretkey, vmttype, quantity	status
setreserveprice	accesskey, secretkey, vmttype, value	price
setmaxprice	accesskey, secretkey, vmttype, value	status
addorder	accesskey, secretkey, vmttype, quantity, bid, ref	price
updateorder	accesskey, secretkey, quantity, ref	price
removeorder	accesskey, secretkey, ref	price
pricehistory	accesskey, secretkey, vmttype, fromdate, todate	price (s)

# OpenStack and Horizon

## OpenStack

An open-source cloud management platform

Developed to control large pools of compute, storage, and networking resources in a data center.

## Horizon

The OpenStack dashboard project

A web based user interface to OpenStack services

Provides administrators with management capabilities

# Extensions for Horizon

To add spot market facilities to OpenStack,

we extended Horizon to be capable of using the services provided by SipaaS.

We added a new panel through which system administrators are capable of enabling spot market support

Maximum and minimum amount of bid price for users,

Number of available VMs for allocation,

Maximum number of VMs a user can request.

We added panels for requesting spot instances and viewing spot market price history

# Requesting spot instances web page

### Request Spot Instances

×

Details \*

Access & Security \*

Post-Creation

Advanced Options

Availability Zone:

nova

Instance Name: \*

test

Flavor: \*

m1.nano

Instance Count: \*

3

Bid Amount: \*

0.0190

Instance Boot Source: \*

Boot from image

Image Name:

cirros-0.3.2-x86\_64-uec (24.0 MB)

Specify the details for launching an instance.

The chart below shows the resources used by this project in relation to the project's quotas.

Flavor Details

Name	m1.nano
VCPUs	1
Root Disk	0 GB
Ephemeral Disk	0 GB
Total Disk	0 GB
RAM	64 MB

Project Limits

Number of Instances0 of 20 Used


Number of VCPUs0 of 20 Used

Total RAM0 of 51,200 MB Used

Cancel

Submit

# Spot Pricing History Webpage



# Evaluation and Validation

We conducted 20-minute experiment with 10 participants (i.e., spot market users).

Participants are selected from a group of professional cloud users who have satisfactory level of knowledge about the spot market.

Participants are divided into two groups of five:

- (i) Group T or truthful bidders and
- (ii) Group C or counterpart bidders who have the freedom to misreport their true private values to maximize their utility

User	Price Value (\$)	Quantity
T1, C1	0.0691	2
T2, C2	0.0092	1
T3, C3	0.0475	1
T4, C4	0.0232	2
T5, C5	0.0184	1


# Utility Function

- Utility function for one time slot instance usage, formulated as below:

$$u(r, b) = \begin{cases} (qv - rp)x, & \text{if } b \geq p \text{ and } r \geq q; \\ 0, & \text{otherwise.} \end{cases}$$

- $r$ : requested number of instances
- $b$ : bid price value
- $q$ : true private quantity
- $v$ : true private price value
- $p$ : spot market price at the time of order submission
- $x$ : A Boolean value describing whether the order is accepted or not, respectively.

# Results



Spot market price fluctuation during the experiment.



# Results (Cont.)

User	Total Cost (\$)	Number of Full Time Slots	Utility Value (\$)
T1	1.2964	16	<b>0.9148</b>
C1	1.8216	17	0.5278
T2	0.0000	0	<b>0.0000</b>
C2	10.0227	18	-9.8571
T3	0.1896	6	0.0954
C3	0.2280	8	<b>0.1520</b>
T4	0.0436	1	<b>0.0030</b>
C4	3.6810	5	-3.4490
T5	0.0000	0	<b>0.0000</b>
C5	0.0738	2	-0.0370

Total cost, the number of full time slots usage, and utility value of experiment participants.

# Take Home Message

- Think problems at a higher level
- Computation problem is not only a technical solution but also an economic mechanism
- We need to learn to formulate the problem, model it, and understand it in real experiments

# Future directions

## Advanced Resource Provisioning Policies

Power Saving, Future load prediction

## Option Trading Strategies

Portfolio, Bulk trading, Selling options

## Game Theoretical Analysis of Cloud Federation

Selfish provider and its own profit, Social welfare

## Customer Diversion in Revenue Management Framework

Bumped by the admission control mechanism

## Revenue Management with Overbooking

Underutilized reserved capacity

## Multi-dimensional Truthful Mechanism Design

Two-dimensionally truthful (bid price and quantity)